

Data Wrangling with Dplyr

Symposium: Using RStudio for Visualization and Analysis of Weed Science Experiments

Maxwel Coura Oliveira, PhD

Department of Agronomy
University of Wisconsin-Madison

December 2019

Outline


- 7 dplyr verbs for data manipulation
 - **select**, **filter**, **group_by**, **summarize**, **mutate**, **count**, **arrange**
- Combining verbs using pipes `%>%`
- 2 tidyr verbs to reshape your data (**spread**, **gather**)

Prerequisites

- Install R and R studio
- See the R basics lesson if you're unfamiliar with R or R studio

Create a new R script

- **File > New File > R script**
- Save it in your project directory
- Look on the top left of the R Studio window to find project directory



What is the tidyverse?

- Packages for data manipulation
- Built for data tables
- Makes data manipulation easier than in base R
- Combine verbs using pipes



Installing packages

```
install.packages()
```

- Input: package name
- Downloads packages from CRAN (Comprehensive R Archive Network)
- Install once per machine

```
install.packages("tidyverse")
```

Loading packages

```
library(packagename)
```

```
library(tidyverse)
```

- **Input:** package name
- Gives R access to functions in the package
- Load packages every time you restart R

Data set: barley yields in Minnesota

- Stored in R as **barley** dataset at lattice package

```
library(lattice)
```

- *Rows*: observations of individual columns
- *Columns*: Variables that describe the experiment
 - yield, variety, year, site



barley data

barley

##	yield	variety	year	site
## 1	27.00000	Manchuria	1931	University Farm
## 2	48.86667	Manchuria	1931	Waseca
## 3	27.43334	Manchuria	1931	Morris
## 4	39.93333	Manchuria	1931	Crookston
## 5	32.96667	Manchuria	1931	Grand Rapids
## 6	28.96667	Manchuria	1931	Duluth
## 7	43.06666	Glabron	1931	University Farm
## 8	55.20000	Glabron	1931	Waseca
## 9	28.76667	Glabron	1931	Morris
## 10	38.13333	Glabron	1931	Crookston
## 11	29.13333	Glabron	1931	Grand Rapids
## 12	29.66667	Glabron	1931	Duluth
## 13	35.13333	Svansota	1931	University Farm
## 14	47.33333	Svansota	1931	Waseca
## 15	25.76667	Svansota	1931	Morris
## 16	40.46667	Svansota	1931	Crookston
## 17	29.66667	Svansota	1931	Grand Rapids
## 18	25.70000	Svansota	1931	Duluth

Import data in tidyverse

- *read_csv()* – loads contents of a CSV file
- *Input*: a file path
- *Output* a “tibble”

Why not read.csv()?

- read_csv() is faster
- Create tibbles
- More reproducible

Data frame	Tibble
Strings to factors	Keeps character
Has row names	No row names
Changes column names	Keeps column names as they are

dplyr verbs

- First argument is always a table

- Tibble or data frame

- Output is a new table

- Doesn't change input data

- Must save the output
using <-

```
new_df <- verb(old_df, ... )
```

OR

```
old_df <- verb(old_df, ... )
```

select()

- Selects columns from a data frame
- Input: data and columns to be kept
- Output: data with only the specified columns

```
select(barley, site)
```

```
##           site
## 1 University Farm
## 2      Waseca
## 3      Morris
## 4    Crookston
## 5   Grand Rapids
## 6      Duluth
## 7 University Farm
## 8      Waseca
## 9      Morris
##10    Crookston
##11   Grand Rapids
##12      Duluth
##13 University Farm
##14      Waseca
##15      Morris
##16    Crookston
##17   Grand Rapids
##18      Duluth
##19 University Farm
##20      Waseca
##21      Morris
##22    Crookston
##23   Grand Rapids
##24      Duluth
##25 University Farm
```

filter()

- Choose rows based on values of a variable
- **Input:** data and a logical expression (returns true/false)
 - `<`, `>`, `>=`, `<=`, `==`, `!=`
- **Output:** data with rows that match the expression

```
filter(barley, site == "Waseca")
```

	yield	variety	year	site
## 1	48.86667	Manchuria	1931	Waseca
## 2	55.20000	Glabron	1931	Waseca
## 3	47.33333	Svansota	1931	Waseca
## 4	50.23333	Velvet	1931	Waseca
## 5	63.83330	Trebi	1931	Waseca
## 6	58.10000	No. 457	1931	Waseca
## 7	65.76670	No. 462	1931	Waseca
## 8	48.56666	Peatland	1931	Waseca
## 9	46.76667	No. 475	1931	Waseca
## 10	58.80000	Wisconsin No. 38	1931	Waseca
## 11	33.46667	Manchuria	1932	Waseca
## 12	37.73333	Glabron	1932	Waseca
## 13	38.50000	Svansota	1932	Waseca
## 14	37.40000	Velvet	1932	Waseca
## 15	49.23330	Trebi	1932	Waseca
## 16	42.20000	No. 457	1932	Waseca
## 17	44.70000	No. 462	1932	Waseca
## 18	36.03333	Peatland	1932	Waseca
## 19	41.26667	No. 475	1932	Waseca
## 20	58.16667	Wisconsin No. 38	1932	Waseca

Pipe operator %>%

- Combine multiple verbs
- **Syntax:** %>% at the end of the line
- Output of the first line becomes input of next line, etc.
- Say it out loud as “then”

```
barley %>%  
  filter(yield > 50) %>%  
  select(site, yield)
```

```
##      site    yield  
## 1 Waseca 55.20000  
## 2 Waseca 50.23333  
## 3 Waseca 63.83330  
## 4 Waseca 58.10000  
## 5 Waseca 65.76670  
## 6 Waseca 58.80000  
## 7 Waseca 58.16667
```

Exercise #1: practice pipes

- Using pipes, subset the **barley** data to include
- **yield** of individuals higher than 40 and lower than 25

mutate()

- Creates a new column
- **Input:** data and the definition of a new column
- `col_name =`
- **Output:** data with a new column

```
barley %>%  
  mutate(yield_kgha = round(yield * 67.25, 0))
```

##	yield	variety	year	site	yield_kgha
## 1	27.00000	Manchuria	1931	University Farm	1816
## 2	48.86667	Manchuria	1931	Waseca	3286
## 3	27.43334	Manchuria	1931	Morris	1845
## 4	39.93333	Manchuria	1931	Crookston	2686
## 5	32.96667	Manchuria	1931	Grand Rapids	2217
## 6	28.96667	Manchuria	1931	Duluth	1948
## 7	43.06666	Glabron	1931	University Farm	2896
## 8	55.20000	Glabron	1931	Waseca	3712
## 9	28.76667	Glabron	1931	Morris	1935
## 10	38.13333	Glabron	1931	Crookston	2564
## 11	29.13333	Glabron	1931	Grand Rapids	1959
## 12	29.66667	Glabron	1931	Duluth	1995
## 13	35.13333	Svansota	1931	University Farm	2363

Exercise 2: data frame challenge

- Create a new data frame from the barley data that meets the following criteria:
 - 1 contains only the **site** and **yield** column and a new column called **yield_lb**
 - 2 **yield_lb** contains values that are yield in lb / 1000 sq. ft
- **Hint:** $1 \text{ bu/acre} = 1.38 \text{ lb/1000 sq. ft}$

Creating a summary table

- `summarize()`
- **Input:** data and a summary statistic
 - Eg: `mean()`
 - `na.rm = TRUE`
- **Output:** a table with the calculated summary statistic

```
barley %>%  
  summarize(mean_yield = mean(yield,  
                              na.rm=TRUE))  
  
##   mean_yield  
## 1    34.42056
```

Creating a grouped summary table

- `group_by()`
- **Input:** data and a variable
- **Output:** a table with the calculated summary statistic for each unique value in the variable

```
barley %>%  
  group_by(site) %>%  
  summarize(mean_yield = mean(yield,  
                              na.rm=TRUE))
```



```
## # A tibble: 6 x 2  
##   site          mean_yield  
##   <fct>          <dbl>  
## 1 Grand Rapids      24.9  
## 2 Duluth            28.0  
## 3 University Farm   32.7  
## 4 Morris            35.4  
## 5 Crookston         37.4  
## 6 Waseca            48.1
```

Removing missing values

- `is.na()`

- missing = **TRUE**
- not missing = **FALSE**

- **Input:** a column

- **Output:** logical vector

- Use it as input to `filter()`

```
barley %>%  
  filter(!is.na(yield)) %>%  
  group_by(variety) %>%  
  summarize(mean_yield = mean(round(yield),  
                                na.rm=TRUE))
```

```
## # A tibble: 10 x 2  
##   variety      mean_yield  
##   <fct>          <dbl>  
## 1 Svansota        30.3  
## 2 No. 462         35.5  
## 3 Manchuria       31.4  
## 4 No. 475         31.8  
## 5 Velvet          32.9  
## 6 Peatland        34.2  
## 7 Glabron         33.3  
## 8 No. 457         35.8  
## 9 Wisconsin No. 38 39.3  
## 10 Trebi         39.6
```

count()

- Count the number of observations
- **Input:**
 - categorical variable
- **sort** = TRUE: sorts the results
- **Output:** a table with a row for each categorical variable and a column called n with counts

```
barley %>%  
  count(site)  
  
## # A tibble: 6 x 2  
##   site          n  
##   <fct>      <int>  
## 1 Grand Rapids 20  
## 2 Duluth       20  
## 3 University Farm 20  
## 4 Morris       20  
## 5 Crookston    20  
## 6 Waseca       20  
Same as  
barley %>%  
  group_by(site) %>%  
  summarize(count=n())  
  
## # A tibble: 6 x 2  
##   site          count  
##   <fct>      <int>  
## 1 Grand Rapids    20  
## 2 Duluth          20  
## 3 University Farm 20  
## 4 Morris          20  
## 5 Crookston       20  
## 6 Waseca          20
```

arrange()

- Order results in ascending order
- **Input:**
 - A variable
 - Use **desc()** to put them in descending order
- **Output:** A table ordered by the values of the input column

```
barley %>%
  group_by(variety) %>%
  arrange(desc(yield))
```

A tibble: 120 x 4

Groups: variety [10]

	yield	variety	year	site
##	<dbl>	<fct>	<fct>	<fct>
## 1	65.8	No. 462	1931	Waseca
## 2	63.8	Trebi	1931	Waseca
## 3	58.8	Wisconsin No. 38	1931	Waseca
## 4	58.2	Wisconsin No. 38	1932	Waseca
## 5	58.1	No. 457	1931	Waseca
## 6	55.2	Glabron	1931	Waseca
## 7	50.2	Velvet	1931	Waseca
## 8	49.9	Wisconsin No. 38	1931	Crookston
## 9	49.2	Trebi	1932	Waseca
## 10	48.9	Manchuria	1931	Waseca
## #	... with 110 more rows			

Exercise 3

- 1 - Use `group_by()` and `summarize()` to find the `mean()`, `min()`, and `max()` `yield` for each `site`.
- 2 - **Bonus:** What is the yield gap between varieties in each site and location?

Reshaping data with tidyr

- The shape of your data affects what you can do with it
- **Example:** compare the mean **yield** of each variety adding a new column (High or Low)



Exercise

- Create a table with columns for **variety** and **mean yield**. Add a logical parameter for mean **yield**, > 35 (High) or < 35 (Low)
- Save to a new object called **barley_nd**


```
barley_nd <- barley %>%  
  select(variety, yield) %>%  
  group_by(variety) %>%  
  summarise(mean_yield = mean(round(yield))) %>%  
  mutate(size = ifelse(mean_yield > 35,  
                        "High", "Low"))
```

barley_nd

```
## # A tibble: 10 x 3  
##   variety      mean_yield size  
##   <fct>      <dbl> <chr>  
## 1 Svansota      30.3 Low  
## 2 No. 462       35.5 High  
## 3 Manchuria     31.4 Low  
## 4 No. 475       31.8 Low  
## 5 Velvet        32.9 Low  
## 6 Peatland       34.2 Low  
## 7 Glabron        33.3 Low  
## 8 No. 457        35.8 High  
## 9 Wisconsin No. 38 39.3 High  
## 10 Trebi         39.6 High
```


Tidy Data

- 1) Each variable has its own column
- 2) Each observation has its own row
- 3) Each value has its own cell
- 4) Each type of observational unit forms a table



Reshaping data with tidyr

- **Spreading**: makes a wider table
 - Unique values in a specified column (key) become variable names
- **Gathering**: makes a longer table
 - Variable names become values in a new column (key)



spread()

- use it when an observation is scattered across multiple rows

- **Input:**

- **data** (a tibble)
- **key** column (values become new column names)
- **value** column (to fill new column variables)

- **Output:** a table with columns for each value of sex

```
barley_spread <- barley_nd %>%  
  spread(key = size,  
         value = mean_yield)
```

```
barley_spread
```

```
## # A tibble: 10 x 3  
##   variety      High    Low  
##   <fct>      <dbl> <dbl>  
## 1 Svansota      NA    30.3  
## 2 No. 462      35.5    NA  
## 3 Manchuria     NA    31.4  
## 4 No. 475      NA    31.8  
## 5 Velvet       NA    32.9  
## 6 Peatland      NA    34.2  
## 7 Glabron      NA    33.3  
## 8 No. 457      35.8    NA  
## 9 Wisconsin No. 38 39.3    NA  
## 10 Trebi       39.6    NA
```

Spread

variety <fctr>	mean_yield <dbl>	size <chr>
Svansota	30.33333	Low
No. 462	35.50000	High
Manchuria	31.41667	Low
No. 475	31.75000	Low
Velvet	32.91667	Low
Peatland	34.25000	Low
Glabron	33.33333	Low
No. 457	35.83333	High
Wisconsin No. 38	39.33333	High
Trebi	39.58333	High

variety <fctr>	High <dbl>	Low <dbl>
Svansota	NA	30.33333
No. 462	35.50000	NA
Manchuria	NA	31.41667
No. 475	NA	31.75000
Velvet	NA	32.91667
Peatland	NA	34.25000
Glabron	NA	33.33333
No. 457	35.83333	NA
Wisconsin No. 38	39.33333	NA
Trebi	39.58333	NA

gather()

- Use when column names are not names of variables, but values of a variable
- Input:
 - **data** (a tibble)
 - **key** column (created from col names)
 - values column (fill the key variable)
 - A range of columns to gather
- **Output**: a long tibble

```
barley_gather <- barley_spread %>%  
  gather(key = size,  
         value = mean_yield, 2:3, na.rm=TRUE)
```

```
barley_gather
```

```
## # A tibble: 10 x 3  
##   variety      size mean_yield  
##   <fct>      <chr>      <dbl>  
## 1 No. 462    High        35.5  
## 2 No. 457    High        35.8  
## 3 Wisconsin No. 38 High        39.3  
## 4 Trebi      High        39.6  
## 5 Svansota    Low         30.3  
## 6 Manchuria   Low         31.4  
## 7 No. 475     Low         31.8  
## 8 Velvet      Low         32.9  
## 9 Peatland    Low         34.2  
## 10 Glabron    Low         33.3
```

Gather

variety <fct>	High <dbl>	Low <dbl>
Svansota	NA	30.33333
No. 462	35.50000	NA
Manchuria	NA	31.41667
No. 475	NA	31.75000
Velvet	NA	32.91667
Peatland	NA	34.25000
Glabron	NA	33.33333
No. 457	35.83333	NA
Wisconsin No. 38	39.33333	NA
Trebi	39.58333	NA

variety <fct>	size <fct>	mean_yield <dbl>
No. 462	High	35.50000
No. 457	High	35.83333
Wisconsin No. 38	High	39.33333
Trebi	High	39.58333
Svansota	Low	30.33333
Manchuria	Low	31.41667
No. 475	Low	31.75000
Velvet	Low	32.91667
Peatland	Low	34.25000
Glabron	Low	33.33333

write_csv

- Writes a data table to a file
- **Input:** a tibble, a file path
- **Output:** a file at the specified file path

```
write_csv(barley_gather,  
          path = "barley_2.csv")
```

Need help

- E-mail: maxoliveira@wisc.edu
- Data Wrangling Cheat Sheet - [Link](#)
- Thanks to Data Camp for sharing slides